DOE Workspace - App Summary
(Repo evidence: README.md, AGENTS.md, b. directives/, c. execution/)

What it is
- A reliability-focused 4-layer AI orchestration system.
- It separates probabilistic orchestration from deterministic scripts, with research feeding both.

Who it's for
- Primary persona: AI operators/automation builders managing SOP-style directives and execution scripts.

What it does
- Stores research in persistent topic folders with indexes and skill arcs.
- Uses markdown directives (SOPs) with goals, inputs, tools, outputs, and edge cases.
- Routes work through an orchestration layer that reads directives and handles failures.
- Runs deterministic Python execution scripts for API calls, data processing, and file operations.
- Supports webhook-triggered workflows mapped by slug in c. execution/webhooks.json.
- Includes domain automations (for example appraisal order workflow and contact intake).

How it works (repo-evidenced architecture)
- Core flow: a. strategy -> b. directives -> orchestration agent -> c. execution scripts.
- Webhook flow: external POST -> webhook slug lookup in c. execution/webhooks.json -> mapped directive/script -> tool actions.
- Integrations evidenced in repo artifacts: Modal deployment references, n8n JSON workflows, Google Sheets, Gmail.
- Not found in repo: execution/modal_webhook.py file referenced in docs.

How to run (minimal getting started)
- 1) cp .env.template .env, then add required keys in .env.
- 2) Install dependencies: pip install anthropic python-dotenv requests modal
- 3) Optional Google setup: add credentials.json, then run a script to generate token.json.
- 4) Deploy webhooks: modal deploy execution/modal_webhook.py (path currently marked Not found in repo).